

High-Level Synthesis with Simulated Annealing Scheduling

Oleg Kelner
Hochschule Hamm-Lippstadt
Lippstadt, Germany
oleg.kelner@stud.hshl.de

Abstract—

I. INTRODUCTION

High-level synthesis (HLS) compiles a behavioral description into a register-transfer-level data path by solving three interdependent sub-problems: *scheduling* assigns each operation of a data-flow graph (DFG) to a control step, *binding* maps each scheduled operation to a concrete functional-unit instance, and *allocation* decides how many instances of each resource type the data path actually contains. All three choices interact through the same resource-usage pattern over time, and jointly they are NP-hard, since the problem subsumes resource-constrained scheduling. Classical HLS tools sidestep this complexity with fast constructive heuristics – As-Soon-As-Possible (ASAP), As-Late-As-Possible (ALAP), and List Scheduling – that build a single schedule greedily, cycle by cycle. These methods are optimal only under restrictive assumptions (e.g. unlimited resources, or a fixed priority order that happens to match the true critical path) and, being purely constructive, cannot revise an earlier decision once made; they therefore get trapped in local optima of the latency–area cost surface.

Simulated annealing (SA) is a natural fit for escaping this trap: unlike greedy list scheduling, it accepts *worsening* moves with a temperature-controlled probability, which lets the search leave local optima that a monotonically improving heuristic cannot escape, while still converging to a stable solution as the temperature is lowered. SA-based data-path synthesis also has direct precedent in VLSI CAD, where it was used to jointly allocate registers, functional units, and interconnect [1]. This places SA in contrast to exact formulations such as Integer Linear Programming (ILP), which can certify optimality but whose solve time grows too quickly with problem size to remain practical on realistic DFGs.

This paper makes the following contributions:

- a joint scheduling + binding + allocation SA formulation driven by a single weighted cost function (Section IV);
- a concrete SA move set and repair mechanism that keeps every candidate schedule feasible after each move (Section V);
- a complexity and empirical runtime-overhead analysis of the scheduler (Section VI);
- an empirical study of how the α/β latency–area weighting reshapes the search landscape around the ASAP-

seeded starting point, on a realistic DFG benchmark (Section VII).

The remainder of the paper first reviews related scheduling and allocation approaches (Section II), then formalizes the problem (Section III), before presenting the SA formulation, algorithm, complexity analysis, and experiments in turn.

II. RELATED WORK

Classical HLS tools schedule with fast constructive heuristics such as ASAP, ALAP, and List Scheduling, or with the more refined Force-Directed Scheduling, which balances concurrency across control steps to reduce resource cost under a fixed latency [2], [3]; exact formulations based on integer linear programming can certify optimal schedules but scale poorly with problem size [4]. Simulated annealing has previously been used to jointly allocate registers, functional units, and interconnect in data-path synthesis [1], and other metaheuristics such as genetic algorithms have likewise been applied to combined HLS scheduling and allocation [5]. This paper follows that SA-based line of work, with one difference: the number of allocated resource instances is itself a move within the same search, rather than being fixed in advance or optimized in a separate pass.

III. PROBLEM STATEMENT

Input. The input is a data-flow graph $G = (V, E)$, where each vertex $v \in V$ is a typed operation (e.g. LOAD, STORE, ADD, SUB, MUL, DIV), and each edge $(u, v) \in E$ denotes a data dependency requiring u to complete before v begins. A resource library specifies, for each resource type, an area cost, a fixed latency in clock cycles, and an available instance count.

Decision variables. A candidate solution assigns to every operation $v \in V$ a start cycle $c(v) \in \mathbb{N}$ and a resource-instance binding $r(v)$, identifying which physical instance of v 's resource type executes it. A solution is *feasible* iff it satisfies:

- 1) **Precedence:** for every edge $(u, v) \in E$, $c(v) \geq c(u) + \text{latency}(u)$;
- 2) **Resource exclusivity:** for any two operations $v_1 \neq v_2$ with $r(v_1) = r(v_2)$, their execution intervals $[c(v), c(v) + \text{latency}(v))$ do not overlap.

Objective. Informally, the goal is to find a feasible (c, r) , together with an instance count per resource type, that minimizes

a weighted combination of total schedule latency and total allocated area; the precise weighted cost function is defined in Section IV.

Complexity. The joint problem is NP-hard, since fixing the resource counts and binding already reduces it to resource-constrained scheduling, itself a well-known NP-hard problem. Here allocation is not fixed but part of the search, so the combined search space – cycle assignment $\times$ binding $\times$ resource count – is strictly larger still. Even for DFGs of modest size, the number of feasible schedules grows far too quickly for exhaustive or exact search to remain practical, which motivates the simulated-annealing-based approach developed in Sections IV and V.

IV. MATHEMATICAL FOUNDATIONS

A. Cost Function

Every candidate solution (c, r, k) from Section III is scored by a single scalar cost,

$$\text{Cost}(c, r, k) = \alpha \cdot L(c) + \beta \cdot A(k), \quad (1)$$

combining two terms. The latency term $L(c) = \max_{v \in V} (c(v) + \text{latency}(v))$ is the schedule length in clock cycles, i.e. the completion time of the last operation. The area term $A(k) = \sum_t k_t \cdot \text{area}(t)$ is the total hardware cost of the allocated resource pool, summing the unit area of each resource type t over its allocated instance count k_t ; note that area depends only on how many instances are allocated, not on how many operations are bound to them. The weights $\alpha, \beta \geq 0$ let a designer trade latency against area: Equation (1) is a linear scalarization of the underlying two-objective (latency, area) problem into the single objective that the search of Section V minimizes, and different choices of α/β correspond to different preferences on that trade-off.

B. Simulated Annealing Formalism

Simulated annealing explores the space of feasible solutions by repeatedly proposing a neighboring solution and deciding whether to move there. Let Cost_{cur} be the cost of the current solution and $\text{Cost}_{\text{next}}$ the cost of a proposed neighbor, and define $\Delta = \text{Cost}_{\text{next}} - \text{Cost}_{\text{cur}}$. The neighbor is accepted with the Metropolis probability

$$P(\Delta, T) = \begin{cases} 1, & \Delta \leq 0, \\ e^{-\Delta/T}, & \Delta > 0, \end{cases} \quad (2)$$

where T is the current *temperature*: improving moves are always accepted, while worsening moves are accepted with a probability that shrinks as Δ grows or T falls. Temperature is lowered geometrically,

$$T_{k+1} = \gamma \cdot T_k, \quad 0 < \gamma < 1, \quad (3)$$

starting from an initial temperature T_0 and stopping once T drops below a minimum threshold $T_{\min}$. At high temperature, Equation (2) accepts almost any move, so the search explores broadly and can escape local optima of the cost surface; as $T \rightarrow T_{\min}$, only improving or mildly worsening moves survive, so the search behaves increasingly like a greedy

local search and settles into a minimum. A slow enough (e.g. logarithmic) cooling schedule is known to converge to a global optimum in the limit, but at a runtime cost that grows accordingly; the finite, geometric schedule of Equation (3) trades that asymptotic guarantee for a bounded number of iterations, which we return to in Section VI.

V. SIMULATED ANNEALING SCHEDULING ALGORITHM

A. State Representation and Initialization

A candidate solution is represented directly as the triple (c, r, k) from Section III: a start cycle $c(v)$ for every operation, a resource-instance binding $r(v)$, and a resource pool $k = (k_t)_t$ giving the number of allocated instances per resource type. The search is initialized with an ASAP schedule computed under the *full* resource pool declared in the configuration – every available instance of every type, not just the minimum needed. Each operation is placed at the earliest cycle its predecessors allow and greedily bound to a free instance, with the repair mechanism of Section V-B resolving any remaining conflicts. This initial solution is feasible by construction and, because the pool is generous, typically already close to latency-optimal; it is far from cost-optimal in another sense, however, since every declared instance counts toward its area whether the schedule actually uses it or not. Section VII-C returns to this starting point in detail, as it turns out to matter for how the search behaves.

B. Neighbor Generation and Repair

At each step, a neighboring solution is produced by applying one randomly chosen move to the current solution:

- 1) **Move forward**: shift a randomly chosen operation one cycle later.
- 2) **Move backward**: shift a randomly chosen operation one cycle earlier (clamped at cycle 0).
- 3) **Rebind**: reassign a randomly chosen operation to a different, compatible instance of its resource type.
- 4) **Add resource**: increase k_t by one for a chosen resource type t .
- 5) **Remove resource**: decrease k_t by one for a chosen resource type t (never below one), automatically rebinding any operations that were using the removed instance.

A move can, on its own, break one of the two feasibility conditions of Section III – for example moving an operation earlier than one of its predecessors, or binding two operations to the same instance at overlapping cycles. Every move is therefore followed by a *repair* step: *dependency repair* recursively pushes an operation's successors forward whenever precedence is violated, and *resource-overlap repair* shifts operations forward whenever two operations sharing an instance now overlap in time. Because repair always restores a feasible solution, the Metropolis test of Equation (2) can be applied directly to the resulting cost, without any additional penalty term for infeasibility.

Algorithm 1 Simulated Annealing Scheduling (SAS)

Require: DFG $G = (V, E)$, resource library, T_0 , $T_{\min}$, γ , iterations per temperature level I

Ensure: best solution found S^*

```
1:  $S \leftarrow \text{ASAP}(G)$  with minimal resource pool
2:  $S^* \leftarrow S$ 
3:  $T \leftarrow T_0$ 
4: while  $T > T_{\min}$  do
5:   for  $i = 1$  to  $I$  do
6:      $S' \leftarrow \text{Repair}(\text{Move}(S))$ 
7:      $\Delta \leftarrow \text{Cost}(S') - \text{Cost}(S)$ 
8:     if  $\Delta \leq 0$  or  $\text{rand}(0, 1) < e^{-\Delta/T}$  then
9:        $S \leftarrow S'$ 
10:      if  $\text{Cost}(S) < \text{Cost}(S^*)$  then
11:         $S^* \leftarrow S$ 
12:      end if
13:    end if
14:  end for
15:   $T \leftarrow \gamma \cdot T$ 
16: end while
17: return  $S^*$ 
```

C. Acceptance, Cooling, and Multi-Run Strategy

Algorithm 1 summarizes the full outer loop. At each temperature level, a fixed number of neighbors are proposed, repaired, and accepted or rejected via Equation (2); the best solution found so far is tracked separately from the current solution, since a worsening move may still be accepted along the way. Once the proposals at a given temperature are exhausted, the temperature is cooled with Equation (3) and the process repeats until T falls below $T_{\min}$. Because the outcome of a single run depends on its random sequence of proposed moves, the scheduler performs several independent runs from different random seeds and reports the best solution found across all of them.

VI. RUNTIME AND OVERHEAD ANALYSIS

A. Theoretical Complexity

Generating a candidate neighbor is cheap: four of the five move types touch a single operation and run in $O(1)$, while rebind is $O(m_t)$ in the number of available instances m_t of the relevant resource type, which is normally small. The repair step dominates instead: dependency repair can, in the worst case, propagate through every remaining successor of the moved operation, and resource-overlap repair can likewise touch every operation sharing that instance, so a single move-and-repair step costs $O(|V| + |E|)$ in the worst case. At a given temperature, I such steps are performed, and the number of temperature levels visited before T drops below $T_{\min}$ is $\lceil \log_\gamma(T_{\min}/T_0) \rceil$. Repeating this for a number of independent runs R , the total worst-case work is

$$O\left(R \cdot \lceil \log_\gamma(T_{\min}/T_0) \rceil \cdot I \cdot (|V| + |E|)\right). \quad (4)$$

Every factor in Equation (4) except the graph size $|V| + |E|$ is a user-controlled parameter of the search, which makes the

runtime-quality trade-off of Section IV-B explicit and tunable, at the cost of no longer offering an asymptotic optimality guarantee.

B. Empirical Overhead

The bound in Equation (4) is a worst case dominated by graph size and the number of accepted iterations; it is not a tight estimate of actual runtime, since a repair step rarely needs to touch the whole graph in practice. Section VII reports the wall-clock overhead measured on the implemented simulator to complement this theoretical bound with concrete numbers.

VII. SIMULATION

The concepts of Sections IV and V are implemented in a C++ simulator built around the model of Section III. This section walks through what the simulator consumes and produces, how its starting point is obtained, and then makes one observation about how the search behaves as the cost function's weights change.

A. Input: Data-Flow Graph and Configuration

The simulator takes two inputs. The DFG is a flat text file, one operation per line, giving an operation's type and its operand identifiers; dependencies are inferred directly from those operand references. The benchmark DFG used throughout this section follows this pattern:

```
t1  = LOAD(D1A1)
t9  = LOAD(D1B1)
t17 = MUL(t1, t9)
t18 = MUL(t2, t10)
t25 = ADD(t17, t18)
...
t511 = ADD(t510, t509)
t515 = DIV(t511, t514)
STORE(t515)
```

The full graph repeats this load-multiply-reduce pattern across several input vectors, combines the resulting partial sums through further addition and subtraction stages, and finishes with a division and a store, for 515 operations over the six operation types LOAD, STORE, ADD, SUB, MUL, and DIV.

Alongside the DFG, a JSON configuration fixes the SA parameters and the resource library:

```
{
  "sa": {
    "initial_temperature": 1,
    "cooling_rate": 0.95,
    "min_temperature": 0.005,
    "iterations_per_temp": 50,
    "runs": 10,
    "alpha": 0.75,
    "beta": 0.25,
    "seed": 6324798
  },
  "resources": {
```

```

"LOAD": {"count": 25, "area": 5, "latency": 1},
"STORE": {"count": 25, "area": 5, "latency": 1},
"ADD": {"count": 25, "area": 5, "latency": 1},
"SUB": {"count": 25, "area": 5, "latency": 1},
"MUL": {"count": 25, "area": 15, "latency": 3},
"DIV": {"count": 25, "area": 40, "latency": 4}
}

```

initial_temperature, min_temperature, and cooling_rate are T_0 , $T_{\min}$, and γ from Equation (3); iterations_per_temp is I and runs is the number of independent restarts R , both from Algorithm 1; alpha and beta are the cost-function weights of Equation (1) – the two values varied in Section VII-D; and seed is a base value from which each of the R runs derives its own, distinct seed (an offset of the run index), so restarts are independent random streams while the whole experiment stays reproducible. Under resources, each entry instantiates one resource type from Section III: count is the maximum number of instances that may ever be allocated, area their unit area cost, and latency their fixed execution latency in cycles.

B. Output: The Resulting Schedule

The simulator’s output is a concrete assignment of every operation to a resource instance and a cycle interval – the (c, r, k) triple of Section V-A, made explicit. An excerpt of the best schedule found for the benchmark DFG: Each

TABLE I
EXCERPT OF THE RESULTING SCHEDULE (BEST SOLUTION FOUND).

Resource	Start	End	Operation
Res_7_LOAD	4	5	t1
Res_7_LOAD	0	1	t2
Res_3_MUL	6	9	t17
Res_1_ADD	9	10	t25

row fixes the start cycle $c(v)$, the binding $r(v)$ – a named instance, not just a resource type – and its cycle interval $[c(v), c(v) + \text{latency}(v))$. The set of distinct instance names appearing in the table, grouped by type, is the allocated pool k . Read cycle by cycle, this table is exactly the control-step schedule a downstream HLS back end would use to drive the datapath’s control logic.

C. From ASAP Baseline to Optimized Schedule

Simulated annealing does not start from an arbitrary or random solution: the initial state S_0 handed to Algorithm 1 is an ASAP schedule computed under the *full* resource pool declared in the configuration, i.e. every count-many instance of every type is available from the start. Each operation is placed at the earliest cycle its predecessors allow and greedily bound to a free instance, with the same repair mechanism of Section V-B resolving any remaining conflicts. Because the pool is generous relative to the parallelism present in the benchmark DFG, this ASAP-derived starting point is already




Fig. 1. Cost vs. temperature, latency-optimizing weighting.

close to latency-optimal – there is little room left to shorten the schedule further. Its area, however, is essentially the configuration’s maximum: every declared instance is allocated whether or not the schedule actually uses it, so a large part of the initial resource pool sits idle.

This asymmetry – an already-good latency paired with a deliberately wasteful area – is not an incidental implementation detail. As Section VII-D shows, it directly shapes how the search behaves once α and β change how costly that idle area is judged to be.

D. Cost-versus-Temperature Landscape

At each temperature level T , a run’s I proposals at that level end with the search in some particular current solution – exactly the solution that is carried forward, unchanged, into the next, cooler level. Rather than summarizing everything that happened during those I proposals, each run records only this single *resulting cost*: one number per run, per temperature level. Aggregating this resulting cost across all R independent runs at a given temperature then yields three quantities: the *best* (the minimum resulting cost across runs), the *worst* (the maximum resulting cost across runs), and the *mean* (the average resulting cost across runs). Every run also uses its own random seed (Section VII-A), so the spread between best and worst reflects genuine run-to-run variation and not a shared random sequence. Plotting these three quantities against temperature, read right to left from T_0 down to $T_{\min}$, is the tool used below to visualize how the search behaves.

Figures 1–3 show this plot for three configurations that differ only in α and β : latency-optimizing ($\alpha = 0.75, \beta = 0.25$), balanced ($\alpha = 0.5, \beta = 0.5$), and area-optimizing ($\alpha = 0.25, \beta = 0.75$), all else held fixed.

The one point these three plots are meant to make concerns the slope near T_0 : how sharply the worst and mean lines rise as the search first moves away from the ASAP-seeded starting point S_0 of Section VII-C. Under the latency-optimizing weighting, S_0 sits in a deep well of the cost surface – it is already close to latency-optimal, so almost any move away from it, even a small one, is charged heavily by the large α , and the rise near T_0 is correspondingly steep. As β grows and α shrinks, the same S_0 is no longer sitting in a well at all: its area is close to the configuration’s maximum, which is now the dominant, heavily weighted term, so S_0 sits closer to a local *peak* than a local minimum. Moves that shed idle




Fig. 2. Cost vs. temperature, balanced weighting.




Fig. 3. Cost vs. temperature, area-optimizing weighting.

resource instances are downhill in almost every direction from there, so the search can lower cost with little resistance, and the initial rise flattens out.

Two caveats apply when reading these three plots side by side. First, the cost function itself changes between them, so the absolute cost values on the y -axis are not directly comparable across figures – only the *shape* of each curve, relative to its own starting cost, carries meaning. Second, that shape is not merely descriptive: it feeds back into the search through the Metropolis criterion of Equation (2), since the acceptance probability depends on Δ in the very units these plots show. Read as a potential-energy landscape, the latency-optimizing configuration presents a narrow, steep-sided well around S_0 ; the area-optimizing configuration instead presents a shallow, almost flat basin that gives way to a steady, near-monotonic fall as the temperature cools and idle resources are progressively removed.

VIII. CONCLUSION

[TODO: Summarize the contribution and restate the headline observation from Section VII-D (how the α/β weighting reshapes the cost-vs-temperature landscape around the ASAP-seeded start) once the three configurations have been fully run. State limitations honestly: a single benchmark DFG (no CHStone/MachSuite generalization test), no quantitative baseline comparison against ASAP or other scheduling heuristics, a fixed geometric cooling schedule (no adaptive/reheating strategy), and no comparison against other metaheuristics (GA, tabu search, PSO) or an exact ILP solution on a small-scale instance for optimality-gap validation. Propose future work: a quan-

titative SA-vs-ASAP evaluation, adaptive cooling, multi-objective Pareto-front search (e.g., NSGA-II-style) instead of scalarized α/β cost, a larger benchmark suite, and integration with a real HLS front end (e.g., Vitis HLS or Bambu) to validate synthesized-hardware results against the simulator's cost model.]

REFERENCES

- [1] S. Devadas and A. Newton, "Algorithms for hardware allocation in data path synthesis," *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, vol. 8, pp. 768–781, 7 1989. [Online]. Available: <http://ieeexplore.ieee.org/document/31534/>
- [2] G. Micheli, "Synthesis and optimization of digital circuits," *Choice Reviews Online*, vol. 32, pp. 32–0950–32–0950, 10 1994.
- [3] P. G. Paulin and J. P. Knight, "Force-directed scheduling for the behavioral synthesis of asic's," *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, vol. 8, pp. 661–679, 1989. [Online]. Available: <https://ieeexplore.ieee.org/document/31522>
- [4] C. T. Hwang, J. H. Lee, and Y. C. Hsu, "A formal approach to the scheduling problem in high level synthesis," *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, vol. 10, pp. 464–475, 1991. [Online]. Available: <https://ieeexplore.ieee.org/document/75629>
- [5] M. Heijligers, L. Cluitmans, and J. Jess, "High-level synthesis scheduling and allocation using genetic algorithms," *Proceedings of ASP-DAC'95/CHDL'95/VLSI'95 with EDA Technofair*, pp. 61–66, 1995. [Online]. Available: <https://ieeexplore.ieee.org/document/486203>